Yanwei (Eric) Feng

yanweife@usc.edu | LinkedIn | GitHub | FinPilot AI (live demo) | Los Angeles, CA

PROFESSIONAL SUMMARY

AI product-oriented builder with hands-on experience developing and deploying LLM-powered systems. Built FinPilot AI, an end-to-end financial reasoning assistant using Qwen 0.6B + LoRA, covering data processing, fine-tuning, evaluation, and AWS deployment. Processed 5,000+ financial instruction samples and designed multi-scenario reasoning flows for risk, strategy, and market analysis. USC M.S. Engineering student (GPA: 3.82) with foundations in optimization, analytics, and statistical modeling.

CORE COMPETENCIES

Product / Program Management | AI & Data Platforms | Cross-functional Coordination | Roadmap & Intake Management | Data Analytics | Stakeholder Communication | Agile Methodologies | LLM Pipelines & Evaluation | Risk & Change Management | Optimization & Statistical Modeling

EXPERIENCE

Shelley Whitman Associates (SWA) — Los Angeles, CA

Project Manager Intern

- **Led cross-functional coordination** across design, engineering, and marketing teams for multiple concurrent client projects
- Managed project timelines and resource allocation using Microsoft Project; surfaced capacity constraints to stakeholders
- Delivered website development project for Mark's Northeastern Furniture on-time by prioritizing intake and resolving bottlenecks
- Produced animated media content for the Salvation Army campaign while balancing competing project deadlines
- Facilitated stakeholder communication to align priorities and resolve scheduling conflicts across teams

PROJECTS

FinPilot AI -- Financial Decision Assistant LIVE

- **Built and deployed** a live LLM-powered financial assistant using Qwen 0.6B + LoRA on AWS EC2
- Processed 5,000+ finance-domain instruction records; added synthetic examples for advanced market and risk analysis
- Designed 5 reasoning modes covering market interpretation, risk analysis, strategy, scenario planning, and general Q&A
- Created product landing page with architecture diagram, demo flow, GitHub link, and live interactive interface

Python, PyTorch, HuggingFace, LoRA/PEFT, FastAPI, AWS EC2, Node.js

LLM Data & Training Pipeline

- Built data processing pipeline for LLM fine-tuning: dataset cleaning, JSONL transformation, 90/10 train-val split, and synthetic data augmentation
- Implemented LoRA fine-tuning with gradient accumulation, evaluation every 100 steps, and checkpointing every 400 steps

Python, HuggingFace Datasets, PyTorch, LoRA, JSONL pipelines

Customer Service & Medical QA Systems

- Built two LLM-based QA systems with multi-step retrieval pipelines; evaluated outputs for accuracy and hallucination reduction
- Applied prompt engineering and domain constraints to improve factual reliability in high-stakes medical and customer service contexts

Python, LLM APIs, Prompt Engineering, Evaluation Frameworks

EDUCATION

University of Southern California (USC) -- M.S. in Industrial & Systems Engineering — Los Angeles, CA
GPA: 3.82 / 4.0 | Coursework: Machine Learning, Optimization & Convex Analysis, Time Series Analysis, Numerical Methods

University of California, San Diego (UCSD) -- B.S. in Applied Mathematics — San Diego, CA
Coursework: Numerical Analysis

TECHNICAL SKILLS

Product & Program: Roadmap management, Intake prioritization, Stakeholder communication, Agile, Microsoft Project

AI/ML: LLM Fine-tuning (LoRA), Prompt Engineering, Model Evaluation, HuggingFace Transformers, Data Pipelines

Analytics: Python, Pandas, Statistical Modeling, Optimization (LP, gradient methods), Time Series Analysis

Tools: AWS EC2, Docker, Git/GitHub, FastAPI, Node.js