

Yanwei Feng

Email: yanweife@usc.edu | **Tel:** +1 951-813-1669 | **GitHub:** <https://github.com/Eric-Fung-Jinwai>

EDUCATION

University of Southern California

Master of Science in Industrial and Systems Engineering

Cumulative GPA: 3.82/4.0

Coursework: Optimization, Machine Learning, ERP Systems, Supply Chain Analytics

Los Angeles, CA

Aug 2024 – June 2026

University of California, San Diego

Bachelor of Science in Applied Mathematics

Cumulative GPA: 3.6/4.0

San Diego, CA

Sep 2022 – Dec 2023

PROFESSIONAL EXPERIENCE

Zhima

AI Engineer

Los Angeles, CA

Dec 2025 - Jun 2026

- Built a production grade multi-agent GTM outreach platform orchestrating four specialized AI agents, reducing orchestration latency by $4.5\times$ through bounded asynchronous execution.
- Developed an evidence-grounded verification framework combining deterministic validation and LLM-based faithfulness evaluation, achieving 97% factual grounding accuracy by validating generated email claims against source website content.
- Built a 176 offline testing regression suite enabling end-to-end validation of the complete multi-agent workflow without external API dependencies.
- Implemented authenticated multi-user infrastructure with session management, tenant isolation, optimistic locking, and server sent events for live campaign progress tracking.

Shelley Whitman Associate

Project Manager

Los Angeles, CA

June 2025 – Sep 2025

- Managed end-to-end development of Mark's Northeast Furniture e-commerce website, tracking 40+ tasks in **OpenProject** and **Jira** to achieve on-time launch.
- Facilitated daily cross-functional departments (web, graphic, illustration) to resolve obstacles, maintain consistency with client requirements and creative vision, and keep deliverables aligned with client goals.
- Assessed risk exposure and optimized resources allocation across three concurrent design initiatives, cutting overtime hours by **15%**.

PROJECT

Citation-Aware Equity Research Assistant

- Built a citation-aware equity research assistant that synthesizes SEC filings, market data, technical indicators, and financial news into structured research reports averaging 12 verifiable citations per report.

- Developed a retrieval-augmented generation (RAG) pipeline indexing 45 SEC filings (10-K, 10-Q, and 8-K) into a ChromaDB vector store containing 16,552 embedded document chunks, enabling citation-grounded equity research from primary source financial evidence.
 - Architected a production level AI research platform using FastAPI and Streamlit with concurrent multi-source retrieval, deterministic confidence validation, and graceful degradation to ensure reliable report generation despite upstream data-source failures.
 - Designed a deterministic confidence scoring and thesis validation framework combining data completeness, signal coherence, retrieval evidence depth, and historical backtesting against SPY benchmarks to generate explainable reliability assessments for AI generated investment reports.
 - Implemented Redis-based multi-layer caching for market data, SEC retrieval, financial news, and report generation, reducing repeat-query latency by 87.5% while eliminating redundant external API requests.
-

SKILLS & ADDITIONAL INFORMATION

- **Programming:** Python, SQL, Java
- **AI/ML:** PyTorch, Transformers, Hugging Face, Scikit-learn, Sentence Transformers, BERT, LSTM
- **LLM Systems:** RAG, QLoRA Fine-Tuning, AI Agents, Vector Databases, Prompt Engineering
- **Backend & Infrastructure:** FastAPI, Redis, Docker, Docker Compose, GitHub Actions, REST APIs
- **Frameworks & Tools:** Streamlit, ChromaDB, Pandas, Git, OpenProject, MS Project, Jira